

The Sky is Falling?

The NEOCP in the Era of LSST

T. WAGG ¹, M. JURIC ¹ AND . . . MORE

¹*Department of Astronomy, University of Washington, Seattle, WA, 98195*

(Received September 15, 2022)

ABSTRACT

[TODO:] Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Keywords: Near-Earth objects, Asteroids, Solar system, Small Solar System bodies, Surveys

1. INTRODUCTION

Near-Earth Objects (NEOs) are objects that may come close enough to the Earth that small perturbations in their orbit lead to a collision course. A simple definition for NEOs is that there are small solar system bodies that have a perihelion less than 1.3 AU (e.g. Jones et al. 2018). Given the threat posed by these objects, it is therefore critical that we catalogue and determine the orbit of NEOs.

The Minor Planet Center maintains a catalogue of known NEOs and their orbits¹, as well as the NEO confirmation page (NEOCP²), which lists potential NEOs that should be prioritised for additional observations by

the NEO follow-up community. An object is only listed on the NEOCP when it has a high probability of being an NEO. This probability is quantified using the **digest2** code, which assigns a score between 0 and 100 and only objects with a score of 65 or more are listed on the page (Keys et al. 2019). Currently, on average around two dozen objects are added to the NEOCP on each night.

The Legacy Survey of Space and Time (LSST) will rapidly increase the rate at which NEOs are discovered and sent to the NEOCP. **[TODO: need more details]** Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gra-

Corresponding author: Tom Wagg
tomjwagg@gmail.com

¹<https://www.minorplanetcenter.net/iau/MPCORB/NEA.txt>

²https://www.minorplanetcenter.net/iau/NEO/toconfirm_tabular.html

vida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Given this increased discovery rate, if LSST operates in the same way as current telescopes that work on NEO discovery, the number of objects submitted to the NEOCP will increase by several orders of magnitude. This will make it difficult for the follow-up community to prioritise particular objects and may lead to some objects being missed. However, many of the objects that LSST observes will be re-observed and discovered by LSST alone, though there will still be a sizeable fraction of objects that require community follow-up. It is therefore important to identify this fraction so that the follow-up community can effectively prioritise observations.

This paper quantifies the impact of LSST on the NEOCP and the fraction of objects that require community follow-up.

Our paper is structured as follows. In Section 2, we describe the methods for simulating LSST observations, calculating `digest2` scores for the NEOCP and our algorithm for predicting whether LSST will detect an object based on a single night of observations. In Section 3, we present our findings for the total number and type of objects submitted to the NEOCP by LSST. Based on these results, in Section 4 we present recommendations for how objects should be submitted to the NEOCP from LSST. We conclude in Section 5.

2. METHOD

In order to make predictions for the NEOCP in the era of LSST, we make simulated observations of a catalogue of solar system objects that takes into account currently known objects. We then use the `digest2` code to calculate NEO scores for each object and use these values to make predictions for the NEOCP. Additionally, we present an algorithm for predicting whether LSST will detect an object given a single night of observations. In the subsections below we explain each of these steps in more details.

2.1. Simulated Observations

In order to investigate the effect of LSST sources on the NEOCP, we use a ‘hybrid’ solar system object catalogue. This catalogue contains both real and synthetic objects, whilst retaining the same overall distributions in position, velocity and absolute magnitude found in the purely synthetic catalogue (see Appendix A).

We perform mock LSST observations using a baseline v2.1 10 year `OpSim` simulation (Bianco et al. 2022). These observations account for both scheduled and unscheduled downtime and simulate the current baseline observing strategy that will be followed by LSST. [TODO: Mario should check this (and the citation)]

2.2. `digest2` Score Calculation

The main criterion for an object to be placed on the NEOCP is for it to receive an NEO score of at least 65. This score from 0 to 100 assesses the probability that the object is an NEO and is calculated using the `digest2` code (Keys et al. 2019). We therefore use the same code to check which of the simulated observations of the hybrid catalogue could be submitted to the NEOCP.

We use `digest2` to calculate the NEO score of each NEO and MBA in the hybrid catalogue that is observed in the first year. We apply three further cuts before considering which submissions are eligible for the NEOCP (needs citation).

1. **Number of observations:** We consider only objects which have at least 2 observations on a given night (though we investigate the effect of making this limit more stringent - see Section. 3.1)
2. **Minimum arc length:** We ensure that each arc is at least 1 arcsecond in length
3. **Maximum time separation:** We set the maximum time between observations to 90 minutes. Thus we only allow tracklets that have at least one pair of observations that occur within 90 minutes of each other

In Figure 1 we show the distribution of `digest2` scores for the NEOs and MBAs in the first year of observations. As expected most NEOs have scores around 100 and most MBAs have scores around 0. However, we highlight that due to the volume of MBAs, the `digest2` score does not perform well in classifying NEOs. In particular, even an object assigned a score of 100 only has around a 3% chance of actually being an NEO. [TODO: not sure if this is right place]

2.3. Estimating LSST Detection Probability

In choosing which objects to submit to the NEOCP from LSST, it is also important to consider whether the object would be detected by LSST without external follow-up. We therefore propose a method of *predicting* whether an observed object will be detected by LSST based on current observations. This requires one to predict both the location and magnitude of the object on subsequent nights, as well as the visit schedule of LSST.




Figure 1. *digest2* scores for all NEOs and MBAs observed in the first year of our simulated LSST observations. The top panel shows normalised histograms, the middle panel shows un-normalised histograms and the final panel shows the ratio of the histogram in the middle panel. Note that the latter two panels are on a logarithmic scale.

In order to determine the location of an object on the following nights one needs to know its orbit. Each object on a given night consists of at least two observations, O , that have the form

$$O = \{\alpha, \delta, t\} \quad (1)$$

where α is the right ascension, δ is the declination and t is the time of observation. One can then determine the proper motion of the object on the sky, $(\dot{\alpha}, \dot{\delta})$, by calculating the change in position over time between the two observations. This determines 4 and the 6 orbital

elements, but the topocentric distance, D and radial velocity, \dot{D} of the object are unconstrained.

We draw a sample of D uniformly in log-space between $[0.1, 10]$ AU and a sample of \dot{D} uniformly between $[-50, 10]$ km s^{-1} to create a grid of (D, \dot{D}) . Combining these with the measured $(\alpha, \delta, \dot{\alpha}, \dot{\delta})$ values (and adding $0.1''$ of scatter to the measured on-sky positions to account for the detector uncertainty), we create a series of possible variant orbits for the object and adjust them to account for the light travel time using THOR (Moeyens et al. 2021). Since we are only interested in NEOs, we mask out any orbits that have a perihelion distance of greater than 1.3 AU. Additionally, we determine the absolute magnitude of the object for each orbit using the HG-system. We convert the apparent magnitudes of the observations in the first night to V-band magnitudes using the LSST colour definitions, assuming a C-type asteroid (Jones et al. 2018). We then use the mean observed V-band magnitude, the assumed distances and a fixed slope parameter of $G = 0.15$ (needs citation) to calculate the current absolute magnitude.

Finally, using *OpenOrb*, we produce ephemerides for the orbits for each visit in the following 14 nights of the predicted schedule that the object could potentially reach based on its observed on-sky motion (Granvik et al. 2009). For each orbit and each field, we check whether the object is within the LSST camera footprint of the field using *rubin.sim*³. We additionally convert the predicted apparent magnitude back to the relevant filter and assess whether the object has an apparent magnitude brighter than the 5σ depth of field. If both of these criteria are met then we assume an observation has been made. We show an example of the predictions produced by this method in Figure 2

Each orbit is defined as producing a detection if it has at least 3 observations on at least 3 nights in a 15 day window, requiring that the minimum arc length is 1 arcsecond and that the maximum separation between observations is 90 minutes. We additionally account for all prior observations in this calculation. For example, if we were determining the probability of detection for an object that is observed on night 10 and has previously been observed on night 8, even if we only predict that it will be observed on night 14 in the next 15 days, this would still be deemed a detection. We estimate the overall probability of the object being detected by LSST as a simple fraction of orbits that produce detections.

3. RESULTS

³https://github.com/lsst/rubin_sim




Figure 2. A demonstration of the LSST detection probability prediction. Black circles (to scale with 2.1° radii) represent a bounding circle of LSST’s camera footprint for each visit in this night’s predicted schedule. For each orbit we plot the object’s predicted location at the start and end of the night and colour them by the assumed distance. In cyan we show true position of the object and outline any visits that produce an observation in cyan also.

3.1. Traffic and purity of the NEOCP using current criteria

In Figure 3, we summarise the effect that LSST submissions would have on the NEOCP if *every* object that met the `digest2` score criterion was submitted. The top panel shows the traffic of the NEOCP, meaning the number of objects that would be submitted to the page, whilst the bottom panel shows the purity, meaning the fraction of objects submitted to the page that are actually NEOs.

The current typical traffic of the NEOCP is on the order of two dozen new submissions per night. We show in the top panel that this traffic would increase by up to 3 orders of magnitude as a result of LSST submissions. Each line corresponding to a different number of minimum nightly observations (each increasing from the original choice of 2 as discussed in Section 2.2). Although the traffic is lower when requiring more observations, even when requiring a minimum of 6 observations the traffic can still reach several thousands of submissions per night, which is far more than the NEOCP is currently equipped to handle.

The purity of the NEOCP is also severely impacted by LSST observations, with the abundance of MBA observations polluting the page as false NEOs. For almost the entire year the page will have a purity below 10%, meaning that only 1 in 10 objects on the page is ac-

tually an NEO, and the value can be as low as 0.2%. As noted in Section 2.2, this is due to the sheer abundance of MBAs in LSST observations. Therefore, given these purity levels a significant amount of follow-up time would be wasted in investigating MBAs masquerading as NEOs.

There are two periodic effects that can be noted in both the traffic and the purity panels. There is a clear seasonal variation over the year as the ecliptic plane moves through the sky. LSST observes from the southern hemisphere and thus in first half of the year when the ecliptic is at lower declinations, more MBAs will be observed [TODO: I think this is backwards...but it definitely peaks in the first half??]. This both increases the traffic and decreases the purity of the NEOCP. In addition to this long term variation, there are also shorter term variations that are most clearly seen in the purity panel. On day 17 and approximately every 30 days after this there is a periodic decrease in the purity of the page. This coincides with the occurrence of a full moon, since LSST has to observe a different part of the sky during this time and this leads to more MBA detections and thus lower purity. [TODO: left this a bit handwavy, need to check with Mario]

Overall, it is clear that if we proceed in the same manner as is currently recommended that the NEOCP will




Figure 3. Traffic (number of objects sent) and purity (fraction of objects sent that are NEOs) of the NEOCP during the first year of LSST if every observation that qualifies for submission is submitted. Each line is plotted using a rolling window of a week to smooth stochastic effects. Different lines correspond to different constraints on the number of observations for a tracklet to be submitted. Nights on which no observations were taken are highlighted with grey areas.

not be able to handle the load. We now consider how we could proceed differently.

3.2. The NEOCP without LSST Detections

LSST will be able to detect and characterise many potential NEOs without any external input. This means that many of the objects submitted to the NEOCP will actual result in wasted follow-up time by the community. We assess the magnitude of this effect using the python package `difi` (Moeyens 2021). This package calculates which objects are detected by LSST, where a detection is defined as occurring when the same object is observed on at least 3 separate nights within a 15 day window,

each with at least 2 observations separated by at most 90 minutes (needs citation).

In Figure 4, we repeat Figure 3, but now only including objects that would not be detected by LSST alone. This represents a best case scenario in which we had foreknowledge of which observed objects would be detected. In this scenario we see that the traffic from LSST rarely exceeds more than 100 objects and is often only a handful. Similarly the purity is much higher, with an average of around 25%. We conclude that if we were to know whether an object would be detected by LSST, the load on the NEOCP would remain at manageable

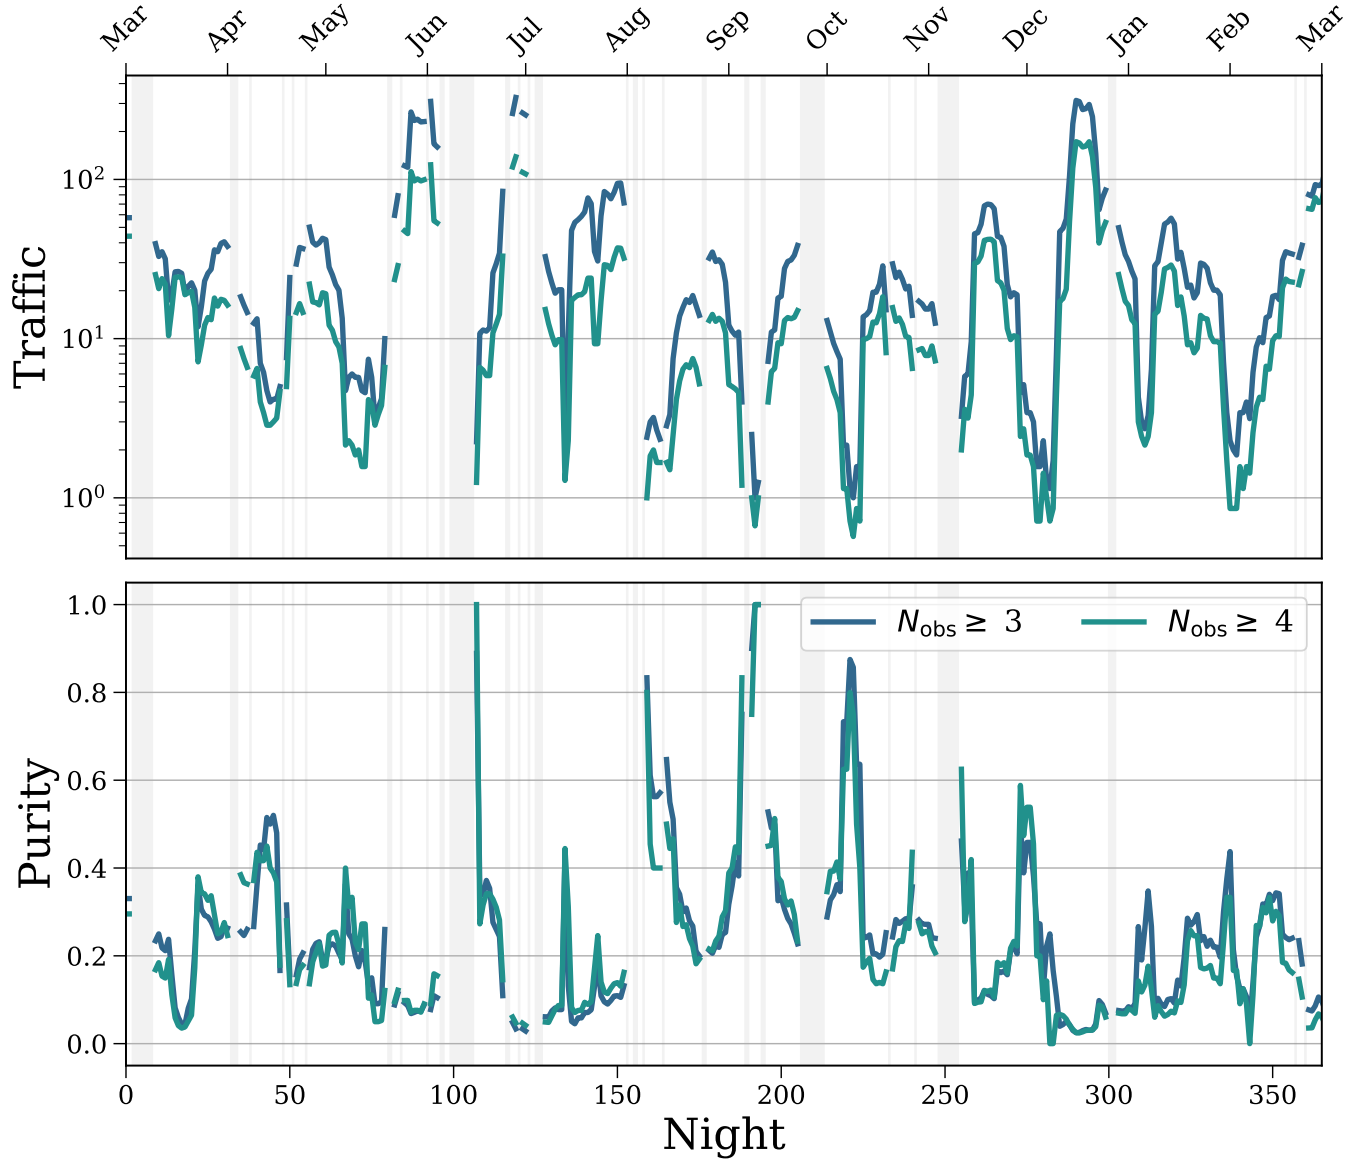


Figure 4. As Figure 3, but only including objects that would not be detected by LSST alone. Note the y-scale of the lower panel is now linear.

levels. For this reason we developed the algorithm that predicts whether an object will be detected by LSST alone (described in Section 2.3) and assess the situation when applying this algorithm in Section 3.3

3.3. Using the LSST detection probability algorithm

We now examine the effect of applying the LSST detection probability algorithm to reduce the load on the NEOCP. Figure 5 shows the results for a year of simulated observations for both the NEOs and MBAs. The histograms show the estimated probability of detection by LSST for each object, split into a population of objects that are *actually* detected within 15 nights in the simulated observations and a population that are not.

On average the algorithm manages to predict the correct outcome approximately 77% of the time.

As one may expect, objects that are detected tend to have higher estimated detection probabilities and vice versa, though there is significant overlap between the distributions. We use a threshold of 0.7 for deciding whether an object will be detected by LSST (and therefore whether to submit it to the NEOCP). We found that this value was ideal for minimising the number of NEOs that are lost and the number of objects sent to the NEOCP, whilst maximising the fraction of those objects that are NEOs.

In the right side of Figure 5, we show contingency matrices for both NEOs and MBAs when applying this

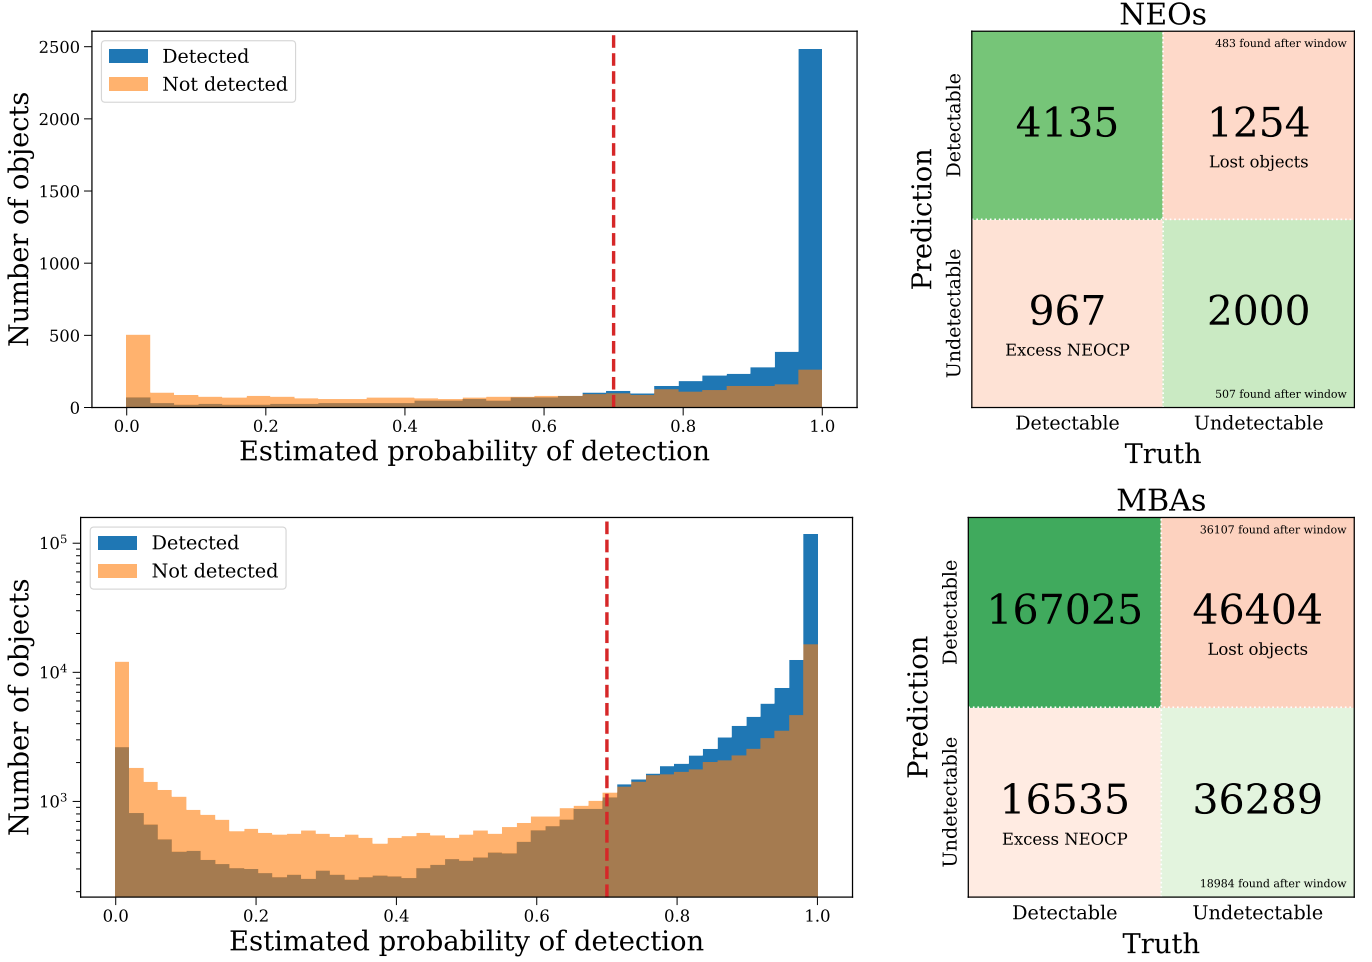


Figure 5. A demonstration of the prediction algorithm described in Section 2.3 using 1 year of simulated observations for both NEOs (top) and MBAs (bottom). **Left:** Probability that an object will be detected within 15 nights, split into a population of objects that are *actually* detected within 15 nights in the simulated observations and a population that are not. The red dashed line indicates our threshold of 0.7 for submission to the NEOCP. Note that the histogram is on a logarithmic scale for the MBAs. **Right:** A contingency matrix representing the algorithm’s ability to predict the detectability of an object.

threshold. The lower left quadrants give a count of the number of objects that will be sent to the NEOCP needlessly as LSST will detect them without follow-up. The upper right quadrants total the number of objects that will be lost as LSST will not detect them within the given detection window but they would also not be sent to the NEOCP. We do additionally annotate the number of objects that would be found later in the year and note this reduces the overall number of lost objects, but increases the number of excess objects in the NEOCP.

Overall, when applying this threshold, we find that, on average, 169 objects would be submitted to the NEOCP per night, but only around 6% (~ 10) of those objects would be NEOs. Moreover, 784 NEOs would remain undetected by LSST and not be sent to the NEOCP for follow-up. This method would reduce the traffic on the NEOCP by well over an order of magnitude but still

has only a small impact on increasing the purity of the sample to acceptable levels.

4. DISCUSSION

As we have shown above, even when applying our LSST detection probability algorithm, the NEOCP would be overwhelmed by submissions from LSST and significantly more MBAs would be submitted than NEOs. In this section we consider recommendations for how to mitigate these issues and ensure that the NEOCP remains functional in the era of LSST.

4.1. Improve *digest2* algorithm

In Figure 1 we highlighted that *digest2* struggles to distinguish between NEOs and MBAs when dealing with the extreme volumes of MBAs that LSST will observe. When using the first year of simulated LSST observations, we find that only $\sim 1.5\%$ of objects with a *digest2*

NEO score of at least 65 (the threshold for submission to the NEOCP) are in fact NEOs.

For this reason, it is important to consider whether there are improvements that can be made to the `digest2` algorithm in distinguishing between NEOs and MBAs. Given the required suppression factor for MBAs and extensive prior work done on this algorithm, it is unlikely that such improvements would be able to eliminate this problem entirely. However, it is possible that there are still particular cases or particular areas of the sky in which the distinction would be improved. For example, the algorithm could take into account the location of the ecliptic at the time of observations.

4.2. *Await reduced MBA background*

The main problem currently stems from the excessive MBA background relative the NEO population. We expect that after the first year of observations LSST will discover approximately 85% of all MBAs down to its flux limit (needs citation). This will strongly suppress the number of MBAs that would be submitted to the NEOCP.

At the same time, we expect a significant fraction of NEOs to be smaller objects that come for infrequent close approaches and so even with new detections we still expect a constant flux of new NEOs (needs citation). In this way the number of NEOs that are submitted to the page would remain roughly constant.

Therefore, once most of the MBAs have been identified, the traffic and purity of NEOCP submissions from LSST would be significantly improved. Indeed, when applying our LSST detection probability algorithm in addition to this, we expect that the resulting traffic and purity would approximate the ideal case described in Section 3.2.

The disadvantage of this approach of course is that any NEOs that will not be re-observed after the first year will not be sent for follow-up and may be lost.

5. SUMMARY & CONCLUSIONS

We present new predictions for the NEOCP in the era of LSST. We perform mock LSST observations and use `digest2` to estimate the number of objects that LSST would send to the NEOCP using the current criteria.

We create a new algorithm for predicting whether an object will be detected by LSST without external follow-up based on a single night of observations in order to reduce the load on the NEOCP. Our main conclusions can be summarised as:

1. **The NEOCP would be overwhelmed by LSST using current submission criteria**

We show that the traffic (number of objects submitted) and purity (fraction of objects that are NEOs) of the NEOCP will be severely impacted by LSST. In particular, the traffic will increase by over two orders of magnitude to between 1000-20000 objects per night, whilst the purity would range from 0.2%-10% (Section 3.1).

2. **Our algorithm would improve the impact on the NEOCP, but not solve the problem**

The application of our LSST detection prediction algorithm would reduce the traffic back to reasonable levels of approximately ~169 objects per night. However the purity still remains at an average of 6% due to the large MBA background (Section 3.3).

3. **NEO follow-up from LSST should not be attempted in the first year**

Due to the vast number of MBA observations in the first year we recommend that follow-up should not be attempted until the second year. At this point around 85% of MBAs will have been identified, whilst the rate of NEO observations will remain roughly constant and so the purity of NEOCP submissions from LSST would return to reasonable levels when using our algorithm (Section 4).

1 [TODO: Acknowledgements]

Software: `digest2` v0.19.2 (Keys et al. 2019), `OpenOrb` (Granvik et al. 2009), `difi` (Moeyens 2021), `THOR` (Moeyens et al. 2021), `Astropy` (Astropy Collaboration et al. 2013, 2018, 2022), `astroML` (VanderPlas et al. 2012, 2014), `scipy` (Virtanen et al. 2020) [TODO: check I didn't miss any]

REFERENCES

- | | |
|---|---|
| <p>Astropy Collaboration, Robitaille, T. P., Tollerud, E. J., et al. 2013, <i>A&A</i>, 558, A33, doi: 10.1051/0004-6361/201322068</p> | <p>Astropy Collaboration, Price-Whelan, A. M., Sipőcz, B. M., et al. 2018, <i>AJ</i>, 156, 123, doi: 10.3847/1538-3881/aabc4f</p> <p>Astropy Collaboration, Price-Whelan, A. M., Lim, P. L., et al. 2022, <i>apj</i>, 935, 167, doi: 10.3847/1538-4357/ac7c74</p> |
|---|---|

- Bianco, F. B., Ivezić, Ž., Jones, R. L., et al. 2022, ApJS, 258, 1, doi: [10.3847/1538-4365/ac3e72](https://doi.org/10.3847/1538-4365/ac3e72)
- Granvik, M., Virtanen, J., Oszkiewicz, D., & Muinonen, K. 2009, M&PS, 44, 1853, doi: [10.1111/j.1945-5100.2009.tb01994.x](https://doi.org/10.1111/j.1945-5100.2009.tb01994.x)
- Grav, T., Jedicke, R., Denneau, L., et al. 2011, PASP, 123, 423, doi: [10.1086/659833](https://doi.org/10.1086/659833)
- Jones, R. L., Slater, C. T., Moeyens, J., et al. 2018, Icarus, 303, 181, doi: [10.1016/j.icarus.2017.11.033](https://doi.org/10.1016/j.icarus.2017.11.033)
- Keys, S., Vereš, P., Payne, M. J., et al. 2019, PASP, 131, 064501, doi: [10.1088/1538-3873/ab1157](https://doi.org/10.1088/1538-3873/ab1157)
- Moeyens, J. 2021, moeyensj/difi: v1.1, v1.1, Zenodo, doi: [10.5281/zenodo.4752719](https://doi.org/10.5281/zenodo.4752719)
- Moeyens, J., Jurić, M., Ford, J., et al. 2021, AJ, 162, 143, doi: [10.3847/1538-3881/ac042b](https://doi.org/10.3847/1538-3881/ac042b)
- VanderPlas, J., Connolly, A. J., Ivezić, Z., & Gray, A. 2012, in Proceedings of Conference on Intelligent Data Understanding (CIDU), 47–54, doi: [10.1109/CIDU.2012.6382200](https://doi.org/10.1109/CIDU.2012.6382200)
- VanderPlas, J., Fouesneau, M., & Taylor, J. 2014, AstroML: Machine learning and data mining in astronomy, Astrophysics Source Code Library, record ascl:1407.018. <http://ascl.net/1407.018>
- Virtanen, P., Gommers, R., Oliphant, T. E., et al. 2020, Nature Methods, 17, 261, doi: [10.1038/s41592-019-0686-2](https://doi.org/10.1038/s41592-019-0686-2)

APPENDIX

A. HYBRID CATALOGUE PIPELINE

Most studies that make predictions for LSST use a synthetic catalogue of solar system objects that doesn't account for prior observations (needs citation). In reality, we have already detected more than a million objects in the solar system and this number will continue to grow until LSST comes online. This means that, current predictions of detection rates will be inflated since a fraction of “new” detections may already be known. Therefore, for this paper we created “hybrid” catalogue that combines a synthetic catalogue with all known observations, whilst keeping the population distributions relatively unchanged.

We created the hybrid catalogue to be dynamic, such that we can run a single pipeline to merge in an updated version of MPCORB as more objects are discovered in the time until LSST comes online. All code to reproduce this hybrid catalogue is open-source and available on GitHub⁴.

A.1. Data preprocessing

For the synthetic catalogue of the solar system we use S3M, the Pan-STARRS Synthetic Solar System Model (S3M Grav et al. 2011). We merge this synthetic catalogue with the latest version of MPCORB⁵, a database of all currently known objects. We use `OpenOrb` (Granvik et al. 2009) to convert both catalogues to Cartesian coordinates and propagate all orbits until the same date.

A.2. Merging algorithm

The general idea for the merging algorithm is to inject each object from MPCORB into S3M, replacing objects that are similar to those injected. An object's similarity is determined based on its position, \vec{x} , velocity, \vec{v} , and absolute magnitude (size), H .

We split each catalogue into bins of absolute magnitude linearly spaced from -2 to 28 and perform the merge algorithm on each bin separately. For each bin we build a K-D trees for both catalogues based on the positions (x, y, z) of objects. For every MPCORB object we query the S3M tree for the nearest 100 objects up to a maximum distance of 0.1 AU, excluding any that have already been matched to a different real object. From these remaining nearest neighbours, we select the S3M object with the closest velocity as the matched object. If there were no remaining neighbours, either because no synthetic objects were nearby or because all nearby objects had already been matched, then we directly add this real object without replacing a synthetic one.

To complete the merging process, we compile the matched object IDs and delete them from S3M. We then add the entirety of MPCORB to the remaining catalogue, resulting in a hybrid catalogue.

A.3. Assessing quality of hybrid catalogue

It is essential that the underlying distributions of the hybrid catalogue do not differ significantly from S3M so that we still accurately reproduce the solar system. In Figure A1, we show the distributions of the absolute magnitude and six orbital elements in both the hybrid catalogue and S3M. It is evident that the distributions are essentially identical.

As a further check, we compared MPCORB to the objects that were removed from S3M, since these should have nearly identical distributions other than MPCORB objects that had no matches. In Figure A2, we show a comparison of the densities for the heliocentric x and y and it is clear that these distributions are left unchanged in the hybrid catalogue.

⁴https://github.com/dirac-institute/hybrid_sso_catalogue/tree/main/hybridcat/hybridcat

⁵<https://minorplanetcenter.net/iau/MPCORB.html>




Figure A1. A comparison of the parameter distributions of S3M (Grav et al. 2011) and the hybrid catalogue we created.




Figure A2. **Left:** A comparison of the density of MPCORB objects with those objects that were matched in S3M by our hybrid catalogue pipeline. **Right:** Residuals between the two plots in the left panel.